

# CSE 368: INTRODUCTION TO AI

Search: Graph search and Local search

Asif Imran, Ph.D.



**UB** University at Buffalo The State University of New York



# Tree Search: Extra Work!


---

- Failure to detect repeated states can cause exponentially more work.



# Graph Search


---



# Graph Search

---

- In BFS, for example, we shouldn't bother expanding the circled nodes (why?)




# Graph Search

---


- Idea: never **expand** a state twice
- How to implement:
  - Tree search + set of expanded states (“closed set”)
  - Expand the search tree node-by-node, but...
  - Before expanding a node, check to make sure its state has never been expanded before
  - If not new, skip it, if new add to closed set
- Important: **store the closed set as a set**, not a list
- Can graph search wreck completeness? Why/why not?
- How about optimality?

# A\* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B }

# A\* Graph Search Gone Wrong?

State space graph



Search tree




Closed set


{ S B }

# A\* Graph Search Gone Wrong?

State space graph




Search tree



Closed set

{ S B C A }


# Consistency of Heuristics




- Main idea: estimated heuristic costs  $\leq$  actual costs
  - Admissibility: heuristic cost  $\leq$  actual cost to goal
$$h(A) \leq \text{actual cost } h^* \text{ from A to G}$$
  - Consistency: heuristic “arc” cost  $\leq$  actual cost for each arc
$$h(A) - h(C) \leq \text{cost}(A \text{ to } C)$$
    - a.k.a. “triangle inequality”:  $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
    - Note: true cost  $h^*$  necessarily satisfies triangle inequality
- Consequences of consistency:
  - The f value along a path never decreases
$$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$
  - A\* graph search is optimal

# A\* Graph Search with Consistent Heuristic

State space graph




Search tree



Closed set

{ S B A C }


# Consistency => non-decreasing f-score



# Optimality of A\* Graph Search

---


- Sketch: consider what A\* does with a consistent heuristic:
  - Fact 1: In tree search, A\* expands nodes in increasing total f value (f-contours)
  - Fact 2: For every state s, nodes that reach s optimally are expanded before nodes that reach s suboptimally
- Result: A\* graph search is optimal



# Optimality

---

- Tree search:
  - A\* is optimal if heuristic is admissible
  - UCS is a special case ( $h = 0$ )
- Graph search:
  - A\* optimal if heuristic is consistent
  - UCS optimal ( $h = 0$  is consistent)
- Consistency implies admissibility
- In general, most natural admissible heuristics tend to be consistent, especially if from relaxed problems



## But...

---


- A\* keeps the entire explored region in memory
- => will run out of space before you get bored waiting for the answer




# Search and Models

---


- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all “in simulation”
  - Your search is only as good as your models...



# Search Gone Wrong?



# Search Gone Wrong?



**Start:** Haugesund, Rogaland, Norway  
**End:** Trondheim, Sør-Trøndelag, Norway  
**Total Distance:** 2713.2 Kilometers  
**Estimated Total Time:** 47 hours, 31 minutes

# Tree Search Pseudo-Code

---

```
function TREE-SEARCH(problem, fringe) return a solution, or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    for child-node in EXPAND(STATE[node], problem) do
      fringe  $\leftarrow$  INSERT(child-node, fringe)
    end
  end
```

# Graph Search Pseudo-Code

---

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
    closed ← an empty set
    fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
    loop do
        if fringe is empty then return failure
        node ← REMOVE-FRONT(fringe)
        if GOAL-TEST(problem, STATE[node]) then return node
        if STATE[node] is not in closed then
            add STATE[node] to closed
            for child-node in EXPAND(STATE[node], problem) do
                fringe ← INSERT(child-node, fringe)
        end
    end
```

# Local Search



---



# Local search algorithms

---

- In many optimization problems, **path** is irrelevant; the goal state **is** the solution
- Then state space = set of “complete” configurations;  
find **configuration satisfying constraints**, e.g., n-queens problem; or, find  
**optimal configuration**, e.g., travelling salesperson problem



- In such cases, can use **iterative improvement** algorithms: keep a single “current” state, try to improve it
- Constant space, suitable for online as well as offline search
- More or less unavoidable if the “state” is yourself (i.e., learning)

# Hill Climbing

---


- Simple, general idea:
  - Start wherever
  - Repeat: move to the best neighboring state
  - If no neighbors better than current, quit



# Heuristic for $n$ -queens problem

---

- Goal:  $n$  queens on board with no **conflicts**, i.e., no queen attacking another
- States:  $n$  queens on board, one per column
- Actions: move a queen in its column
- Heuristic value function: number of conflicts




# Hill-climbing algorithm

---

```
function HILL-CLIMBING(problem) returns a state
    current ← make-node(problem.initial-state)
    loop do
        neighbor ← a highest-valued successor of current
        if neighbor.value ≤ current.value then
            return current.state
        current ← neighbor
```

*“Like climbing Everest in thick fog with amnesia”*

# Global and local maxima



## Random restarts

- find global optimum
- duh

## Random sideways moves

- Escape from shoulders
- Loop forever on flat local maxima


# Hill-climbing on the 8-queens problem

- No sideways moves:

- Succeeds w/ prob. 0.14
- Average number of moves per trial:
  - 4 when succeeding, 3 when getting stuck
- Expected total number of moves needed:
  - $3(1-p)/p + 4 = \sim 22$  moves

- Allowing 100 sideways moves:

- Succeeds w/ prob. 0.94
- Average number of moves per trial:
  - 21 when succeeding, 65 when getting stuck
- Expected total number of moves needed:
  - $65(1-p)/p + 21 = \sim 25$  moves



**Moral: algorithms with knobs to twiddle are irritating**


# Variants of Hill Climbing

- **Stochastic Hill Climbing** selects at random from the uphill moves. The probability of selection varies with the steepness of the uphill move. In fact it selects a random state from the available better states. This usually converges slower than steepest ascent, but in some state landscapes it finds better landscapes
- **First-Choice Climbing** implements the above one by generating successors randomly until a better one (i.e. the first found better state) is found.
- **Random-restart hill climbing** searches from randomly generated initial moves until the goal state is reached.



# Gradient descent

Hill-climbing in continuous spaces



- Gradient vector


$$\nabla J(\underline{\theta}) = \left[ \frac{\partial J(\underline{\theta})}{\partial \theta_0} \quad \frac{\partial J(\underline{\theta})}{\partial \theta_1} \quad \dots \right]$$

- Indicates direction of steepest ascent  
(negative = steepest descent)

# Gradient descent


Hill-climbing in continuous spaces

Gradient = the most direct direction up-hill in the objective (cost) function, so its negative minimizes the cost function.



\* Assume we have some cost-function:  $J(x_1, x_2, \dots, x_n)$  and we want minimize over continuous variables  $x_1, x_2, \dots, x_n$

1. Compute the *gradient*:  $\frac{\partial}{\partial x_i} J(x_1, \dots, x_n) \quad \forall i$



2. Take a small step downhill in the direction of the gradient:

$$x'_i = x_i - \lambda \frac{\partial}{\partial x_i} J(x_1, \dots, x_n)$$

3. Check if  $J(x'_1, \dots, x'_n) < J(x_1, \dots, x_n)$

(or, Armijo rule, etc.)

4. If true then accept move, if not “reject”.

(decrease step size, etc.)

5. Repeat.

# Local beam search


---

- Basic idea:
  - $K$  copies of a local search algorithm, initialized randomly
  - For each iteration
    - Generate ALL successors from  $K$  current states
    - Choose best  $K$  of these to be the new current states

Or,  $K$  chosen randomly with  
a bias towards good ones

# Beam search example ( $K=4$ )


---





# Local beam search

---

- Why is this different from  $K$  local searches in parallel?
  - The searches **communicate**! “Come over here, the grass is greener!”
- What other well-known algorithm does this remind you of?
  - Evolution!



# Genetic algorithms



- Genetic algorithms use a natural selection metaphor
  - Resample  $K$  individuals at each step (selection) weighted by fitness function
  - Combine by pairwise crossover operators, plus mutation to give variety